


AntibioticsANOVA

A student has conducted an experiment evaluating the effect of three different antibiotics — **ab1**, **ab2**, and **ab3** — on the growth of *E. coli* in lab cultures. The results are presented in graphical format below:



We also include charts summarizing the data: one for the overall dataset, one for each type of treatment, and one for the control. Below is the overall summary:

Factors are dropped from the summary

	N	Mean	SD	Min	Q1	Median	Q3	Max
growth	16	12.64	7.25	2.88	5.8	14.03	19.98	21.67

Below is the summary for **ab1**:

Factors are dropped from the summary

	N	Mean	SD	Min	Q1	Median	Q3	Max
growth	4	4.11	1.89	2.88	3.06	3.31	5.16	6.93

Below is the summary for **ab2**:

Factors are dropped from the summary

	N	Mean	SD	Min	Q1	Median	Q3	Max
growth	4	17.52	2.63	14.78	15.71	17.13	19.34	21.04

Below is the summary for **ab3**:

Factors are dropped from the summary

	N	Mean	SD	Min	Q1	Median	Q3	Max
growth	4	8.34	3.64	4.67	5.82	7.69	10.85	13.28

Below is the summary for the **control**:

Factors are dropped from the summary

	N	Mean	SD	Min	Q1	Median	Q3	Max
growth	4	20.61	0.91	19.46	19.98	20.65	21.23	21.67

By examining the data, we might expect **ab1** and **ab3** to be significantly different from the control group; **ab2** does not appear significantly different. Thus we should expect that our ANOVA will return that the means of at least two groups differ. Let μ_0 represent the mean growth of the control group, μ_1 represent the mean of the group treated by **ab1**, μ_2 represent the mean of the group treated by **ab2**, and μ_3 represent the mean of the group treated by **ab3**. Then we have our hypotheses:

Null Hypothesis (H_0): There is not a significant difference in the means of the data: $\mu_0 = \mu_1 = \mu_2 = \mu_3$.

Alternate Hypothesis (H_1): At least one of the antibiotics **ab1**, **ab2**, **ab3** has a mean significantly different from the control group. That is, for some $i \in \{1, 2, 3\}$ we have $\mu_i \neq \mu_0$.

To test these hypotheses we ran an ANOVA-design linear model, and compared the results using a likelihood ratio test. Let G denote growth, let X_i denote treatment by **abi** for $i \in \{1, 2, 3\}$, and let ϵ denote a normally distributed error term. We fit a null model of the form $G = \beta_0$ to this data, optimized for the value of β_0 which minimized the negative log likelihood, and then stored that negative log likelihood value.

For our alternate model, we fit a generalized linear model of the form

$$G = g(X_1, X_2, X_3) = \beta_0 + \beta_1 \cdot X_1 + \beta_2 \cdot X_2 + \beta_3 \cdot X_3 + \epsilon$$

Note that this is a function $g : \mathbb{R}^3 \rightarrow \mathbb{R}$ so we defined our domain as a subset of \mathbb{R}^3 . X_1 took on the values (0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0), X_2 took on the values (0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0), and X_3 took on the values (0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1). The values of 1 in these vectors correspond to the values of $G = g(X_1, X_2, X_3)$ for which the treatments **ab1**, **ab2**, and **ab3** were affected, respectively.

We optimized the model $g(X_1, X_2, X_3)$ for values of the β_i which minimized the negative log likelihood and then stored this value. Then we found D , which is twice the difference of our two stored negative log likelihood values. Using a one-tailed χ^2 distribution with 3 degrees of freedom — the difference in the amount of parameters between the models — we found the p-value associated with a χ -score of D . The p-value was approximately $2.97 \cdot 10^{-8}$.

As the p-value is smaller than 0.05, we reject the null hypothesis and conclude that there is a significant difference between at least two of the means associated with different treatment types.